

# Modern Web Search

Architecture, Algorithms, and Analysis



Image Sources: [www.chegg.com](http://www.chegg.com), [yorko.github.io](http://yorko.github.io), [linkurious.com](http://linkurious.com), [bytebytogo.com](http://bytebytogo.com), [fri-datasience.github.io](http://fri-datasience.github.io), [analytics4all.org](http://analytics4all.org), [medium.datadriveninvestor.com](http://medium.datadriveninvestor.com), [en.wikipedia.org](http://en.wikipedia.org)

# Introduction Part I: Foundations

Scarcity vs. Abundance & Evaluation Metrics

# The Information Retrieval Problem

## Scarcity vs. Abundance

Classic IR (e.g., library systems) dealt with **scarcity**: finding the one relevant document in a pile.

Web Search deals with **abundance**: finding the *best* document among millions of relevant ones.

**Goal:** Satisfy user intent (Navigational, Informational, Transactional) with high precision.

### ⌚Precision

Fraction of retrieved docs that are relevant. Critical for web search (users rarely look past page 1).

### Recall

Fraction of relevant docs that are retrieved. Less critical on the web due to redundancy.

# The Paradigm Shift

---

## Classic IR (Library Science)

Context: A lawyer searching Case Law.

Problem: **Scarcity**. There might be only one relevant precedent.

Goal: Find everything. Missing a document is expensive.

Metric: High Recall.

## Web Search

Context: "Chocolate Cake Recipe".

Problem: **Abundance**. There are 5 million relevant recipes.

Goal: Find the *best* one immediately.

Metric: High Precision (specifically at the top).

# Evaluation Metrics

---

## Basic Metrics

$$\text{Precision} = \text{Rel\_Retrieved} / \text{Total\_Retrieved}$$

$$\text{Recall} = \text{Rel\_Retrieved} / \text{Total\_Relevant}$$

Recall is impossible to calculate on the web (we don't know Total\_Relevant).

## Rank-Aware Metrics

P@K: Precision at top K results (e.g., P@10). "Did I get good stuff on the first page?"

MAP: Mean Average Precision. Rewards placing relevant items higher.

MRR: Mean Reciprocal Rank. For factoid queries (1/Rank of first correct answer).

## Partial Relevance

NDCG: Normalized Discounted Cumulative Gain.

- Handles non-binary scores (0=Spam, 5=Perfect).
- Penalizes good docs if they appear low in the list.

# Exercise: Calculating Metrics

---

## Scenario

Query: "Jaguar" (Intent: Animal)

Results:

1. Doc A (Car)

2. Doc B (Animal)

3. Doc C (Animal)

4. Doc D (OS)

5. Doc E (Animal)

.

## Calculations

P@1: 0/1 = 0.0

P@3: 2/3 = 0.66

Average Precision (AP):

$(P@2 + P@3 + P@5) / 3 = (0.5 + 0.66 + 0.60) / 3 = 0.58$

# Introduction Part II: Web Search Entities

IPs, Domains & DNS

HTML, CSS & JS

# The Address Book

---

## IP Address

142.250.185.78

The "GPS Coordinate". Machines use this.

Pros: Direct, Fast.

Cons: Hard to remember, can change.

## Domain Name

[www.google.com](http://www.google.com)

The "Business Label". Humans use this.

DNS: The phonebook that translates Domain → IP.

Virtual Hosting: One IP can host 1000 domains.

# Domains & DNS

---

## TLD

Top Level Domain. `.com`,  
`.org`, `.gov`. Hints at the  
nature of the site.

## Subdomain

`blog.site.com`. Often treated  
as a separate "site" by search  
engines depending on context.

## DNS

Domain Name System. The  
phonebook of the internet.  
Resolves `google.com` to IP  
`142.250.190.46`.

**Impact on Crawling:** A crawler must resolve DNS for every new host. Efficient DNS caching is critical for performance.

# Anatomy of a URL

Uniform Resource Locators (URLs) are the addresses of the web. Understanding their structure is vital for crawling.


`https://` **Protocol:** How to transfer (HTTP/S).

`www.example.com` **Host/Domain:** The server address.

`/path/to/page` **Path:** Resource location on server.

`?id=123` **Query:** Parameters for dynamic pages.

`#section` **Fragment:** Anchor within the page (not sent to server).



# Anatomy of a URL – example

---



```
https://maps.google.com:443/search?q=paris#center
```

- https:// Protocol (How to talk)
- google.com Domain + TLD (Entity)
- /search Path (File location)
- #center Fragment (Browser scroll pos - NOT sent to server)
- maps Subdomain (Specific service)
- :443 Port (Door number)
- ?q=paris Query (Parameters)

# Semi-Structured Data

The web is **semi-structured**. It has tags (HTML), but the text within is free-form.



## Unstructured

Plain text, images, video.

Hardest to process. Requires  
NLP and Computer Vision.



## Semi-Structured

HTML, XML, JSON. Tags provide clues, but content varies. The standard for web pages.



## Structured

Databases, CSVs, Schema.org.  
Highly organized. Ideal for "rich snippets" (prices, ratings).

# The Language of the Web

---

## HTML (Structure)

The Nouns & Verbs.

Title Link

This is what crawlers care about most.

## CSS (Presentation)

The Adjectives.

.price { color: green; }

Crawlers use this to detect hidden text (spam).

## JS (Behavior)

The Actions.


btn.onclick = () => {  
 fetch('/data') }

Challenge: Crawlers must execute JS to see content on modern sites.

# HTML: The Language of Search

Search engines don't "see" pages like humans. They parse the DOM (Document Object Model).

- **Tags:** </div>, provide semantic weight.
- **Links:** are the edges of the web graph.
- **Meta:** provides snippets.



# The JavaScript Challenge

## Client-Side Rendering

Modern sites (React, Vue) send empty HTML shells.

Content loads via JS.

**Simple Crawlers** see nothing.

**Modern Crawlers** must be "Headless Browsers" (e.g.,  
Puppeteer) to execute JS, render the DOM, and *then*  
extract links.

# **Web Search - Part I**


## **Flow & Architecture**

From Classic IR to Distributed Web Search

# Classic IR Pipeline

The standard process for static document collections  
(e.g., legal corpuses).

1. **Tokenization:** Splitting text into words.
2. **Normalization:** Lowercasing, removing punctuation.
3. **Stopword Removal:** Removing "the", "and" & "is".
4. **Stemming/Lemmatization:** "Running" -> "Run".
5. **Indexing:** Creating the Inverted Index.




# Web Search Architecture

## The Big Picture

The system is cyclic. The crawler feeds the indexer, which feeds the query processor. Analytics feed back into crawling priority.

- **URL Frontier:** The "To-Do" list of links.
- **Fetcher:** Downloads pages (async/multithreaded).
- **Parser:** Extracts links and text.
- **Inverted Index:** Maps words to document IDs.



# The Infinite Web

## Why is it hard to measure?

- **Dynamic Content:** Calendar pages that generate infinite URLs (/day/1, /day/2...).
- **The Deep Web:** Content behind login screens or unlinked databases.
- **Soft 404s:** Pages that say "Error" but return a "200 OK" status.

~60 Trillion

Known Pages (estimated)

Crawlers must prioritize. You cannot index everything.

# Why Distributed?

No single machine can hold the web. We scale horizontally.

**Sharding:** Splitting the index by Document ID (each machine holds a subset of docs) or by Term (each machine holds a subset of words).

**MapReduce:** The paradigm invented by Google to process these massive datasets (e.g., counting links, building the index).

## Challenges

- Fault Tolerance (Machines die).
- Consistency (Index updates).
- Latency (Querying 1000s of machines in <200ms).

# Understanding Intent

Queries are not just keywords; they are goals.

## Navigational

"Facebook", "United Airlines"

User wants to go to a specific site. **High Precision** required on the #1 result.

## Informational

"Who is the CEO of Apple?",

"Symptoms of flu"

User wants to learn. Requires authoritative content.

## Transactional


"Buy iPhone 15", "Download WinZip"

User wants to do something. Requires commerce/action links.

# **Web Search - Part II: Expanding Classic IR to Web Search**

The flow, robots.txt & Scraping

# The Spider's Loop



# Politeness & Robots.txt

---

## DoS vs DDoS

DoS: Single crawler hammering a server. Crash caused by speed.

DDoS: Distributed crawler (1000 machines) hitting a server at once.

Solution: Crawl-delay and Robots.txt.

```
User-agent: * Disallow: /admin/ Crawl-delay: 5  
User-agent: BadBot Disallow: /
```

# Web Search Additions

Web search requires distinct components on top of Classic IR.

## The Crawler

Documents don't just "exist".  
They must be actively  
discovered and fetched from  
remote servers.

## Link Graph

Maintains the map of links  
between pages to compute  
authority scores (PageRank).

## Deduping

Detecting mirrors, scraping, and  
versioning. ~30% of the web is  
duplicate content.

# Scraping vs. Crawling

## Crawling (Discovery)

**Goal:** Discovery & Indexing.

**Scope:** Broad / The whole web.

**Output:** A list of URLs and cached text.

**Example:** Googlebot, Bingbot.

**Tool:** requests, Scrapy.

## Scraping (Extraction)

**Goal:** Extraction.

**Scope:** Narrow / Specific sites.

**Output:** Structured data (prices, names).

**Example:** Price comparison, Lead gen.

**Tool:** BeautifulSoup (Static), Selenium (Dynamic).

# Static Scraping

## The Mechanics

**Pros:** Very fast, low CPU usage.

**Cons:** Cannot see content loaded by JavaScript.

**Libraries:** requests, BeautifulSoup.

```
import requests from bs4
import BeautifulSoup

# 1. Fetch
raw HTML response = requests.get("https://example.com")

# 2. Parse
soup = BeautifulSoup(response.text, 'html.parser')

# 3. Extract specific element (e.g. $19.99 price)
price = soup.find('span', class_='price').text
print(price) # Output: $19.99
```

# Dynamic Scraping

## The Mechanics

**Pros:** WYSIWYG. Can handle infinite scroll and SPAs (React/Vue).

**Cons:** Slow. Launches a full browser instance.

Libraries: Selenium, Puppeteer.

```
from selenium import webdriver
from selenium.webdriver.common.by import By
import time

# 1. Launch Browser
driver = webdriver.Chrome()
driver.get("https://example.com/dynamic")

# 2. Wait for JS to render
time.sleep(2)

# 3. Extract content generated by JS
price = driver.find_element(By.CLASS_NAME, "price").text
print(price)
driver.quit()
```

# Web Search – Part III: Crawling

Crawling, Near-Duplicate Detection

# The "Seed" Problem

You can't crawl the web if you don't know where to start. The web graph must be traversed from known entry points.

## What makes a good seed?

- **High Out-Degree:** Links to many other pages.
- **Trustworthy:** Not spam.
- **Centrality:** Stable and long-standing.
- **Topical Coverage:** Represents the domain well.



# Strategies for Finding Seeds

## Directories

DMOZ (archived), BOTW, and niche directories are human-curated lists of quality sites.

## Institutions

University homepages (.edu) and Government portals (.gov) are high-trust hubs.

## Inverse Links

Use existing indices (like Ahrefs/Majestic) to find who links to known authorities.

## #Social

Highly shared URLs on Twitter/Reddit often point to fresh, relevant content.

## Curated Lists

"Best of" blog posts and Wikipedia "External Links" sections.

## HITS

Run HITS on a small set to identify "Hubs", then use those Hubs as seeds for a larger crawl.

# Crawling Strategy: BFS

## Breadth-First Search

Goal: General Coverage.

Logic: Use a FIFO Queue. Crawl level  $d$  before level  $d+1$ .

Why? Prevents getting stuck in "rabbit holes" (infinite calendars or archives) on a single site.

### Example

- Seed: Home (d=0)
- Q: [Sports, News] (d=1)
- Visit Sports: Adds [Article A, Article B]
- Q: [News, Article A, Article B]
- Visit News: (Crucial: We visit News *before* Article A).

# Code: Simple BFS Crawler

```
def simple_BFS_crawler(start_url, max_pages=10):
    visited = set()
    queue = deque([start_url])
    domain = urlparse(start_url).netloc

    while queue and len(visited) < max_pages:
        url = queue.popleft()
        if url in visited:    continue
        try:
            # 1. Fetch
            response = requests.get(url, timeout=5)

            # 2. Parse
            soup = BeautifulSoup(response.text, 'html.parser')
            visited.add(url) print(f"Crawled: {url}")


            # 3. Extract Links
            for link in soup.find_all('a', href=True):
                abs_url = urljoin(url, link['href'])
                # Filter: Internal links only
                if urlparse(abs_url).netloc == domain:
                    queue.append(abs_url)
        except Exception as e:
            print(f"Error: {e}")
```

# Near-Duplicate Detection

## The Problem

Approximately 30% of the web consists of near-duplicates (mirrors, revisions, spam). Exact match detection fails on dynamic timestamps or minor edits.

Exact hash matching (MD5) fails if one-byte changes (e.g., dynamic timestamp).



## The Solution: Shingling

**Shingles:** Convert text into sets of N-grams (e.g., "a rose is a rose").

**Jaccard Coefficient:** Measure similarity by dividing the intersection of shingle sets by their union.

**MinHashing:** A technique to approximate Jaccard similarity efficiently at scale.

# Near-Duplicate Detection

## Jaccard Similarity

Measures overlap between two sets.

Set A: {apple, banana}

Set B: {apple, cherry}

Intersection: {apple} (Size 1)

Union: {apple, banana, cherry} (Size 3)

Jaccard = 1 / 3 ≈ 0.33

**Pros:** Exact.

**Cons:** Storing sets is expensive.

## MinHash (Approximation)

Goal: Estimate Jaccard with constant space.

1. Create random permutations of vocabulary.
2. For each doc, find the *\*first\** word in the permutation that exists in the doc.
3. If MinHashes match, docs are likely similar.

Insight:  $P(h(A) = h(B)) = \text{Jaccard}(A, B)$ .

# Web Search – Part IV: Link Analysis

Weighted Graphs, PageRank, HITS & Quality

# Weighted Graphs in Search

---

## Dijkstra (Shortest Path)

Use Case: Latency Minimization.

Weights: Ping time (ms).

Method: Priority Queue. Always expand the "cheapest" node ( $d[v]$ ).

Result: Find the fastest server to crawl.

## Prim (MST)

Use Case: Network Backbone.

Weights: Cost of laying cable/bandwidth.

Method: Grow a single tree. Always add the cheapest edge connecting Tree to Non-Tree.


Result: Connect all data centers with min cost.

# The Web as a Graph

The web is a directed graph  $G=(V, E)$ .

- **Vertices ( $V$ ):** Web Pages.
- **Edges ( $E$ ):** Hyperlinks.

Links are not just transport; they are **votes**. A link from A to B is an endorsement of B by A.



# The Bow-Tie Structure

The web is not one big blob. It has distinct regions (Broder et al., 2000).

## IN

New pages. Can  
reach SCC but cannot  
be reached from it.

## SCC

Strongly Connected  
Component. The  
core. Everyone can  
reach everyone.

## OUT

Corporate sites.  
Reached from SCC,  
but link nowhere.


## Tendrils

Isolated islands and  
tubes connecting IN  
to OUT directly.

# The Bow-Tie Structure

The Web is not a purely connected graph. It forms a distinct macroscopic structure:

- **SCC (28%)**: Strongly Connected Component. You can navigate from any page to any other.
- **IN (22%)**: Pages that link to the SCC but cannot be reached from it (New pages).
- **OUT (22%)**: Pages reached from SCC but do not link back (Corporate sites).
- **Tendrils**: Disconnected components and tubes.



# Finding the SCC (Kosaraju)

## The Algorithm

1. Run DFS on Graph  $G$ . Record "Finish Times".
2. Compute Transpose Graph  $G^T$  (Flip all arrows).
3. Run DFS on  $G^T$ , processing nodes in decreasing order of Finish Times.
4. Each resulting tree is an SCC.

### Why Transpose?

By flipping edges, we trap the DFS inside the SCC. It can't "leak out" to other components because the bridges connecting components have been reversed.

# PageRank (PR): The Random Surfer

Imagine a surfer clicking random links forever.


**PageRank (PR)** is the probability that the surfer is on a specific page at any given time.

More incoming links = Higher probability.

Links from high-probability pages = More weight.

Damping factor ( $d$ ): Probability of clicking (usually around 0.85).

Teleportation ( $1-d$ ): Prevents getting stuck in dead ends (Sink nodes).




# Page Rank - The Formula

Recursive definition with Damping Factor.

$$PR(A) = (1-d) + d * ( PR(T_1)/C(T_1) + \dots + PR(T_n)/C(T_n) )$$

- **PR(A)**: PageRank of page A.
- **d**: Damping factor (~0.85). Probability of clicking.
- **(1-d)**: Teleportation. Probability of jumping to a random page (solves dead ends/sinks).
- **C(T<sub>n</sub>)**: Count of outbound links from page T<sub>n</sub>.



# PageRank Calculation

## Scenario

Graph: Node A links to Node B (A  $\rightarrow$  B).

Damping (d): 0.85 (Prob. of clicking).

Teleport (1-d): 0.15 (Prob. of jumping).

$$PR(A) = (1-d) + d * \sum (PR(T)/C(T))$$

Step 1: Calculate PR(A)  
Incoming Links: None.

$$PR(A) = (1-d) = 0.15$$

Step 2: Calculate PR(B)  
Incoming Links: A.

$$\begin{aligned} PR(B) &= (1-d) + d * ( PR(A) / \text{OutDegree}(A) ) = \\ &= 0.15 + 0.85 * ( 0.15 / 1 ) = \\ &= 0.15 + 0.1275 = \\ &= 0.2775 \end{aligned}$$

Result: B is more important than A.

# PageRank

---

## The Random Surfer

PR represents the probability of a random user landing on a page.

## HITS (Alternative)

Hubs: Good lists of links.

Authorities: Good content.

Mutually reinforcing relationship (Recursive).

# HITS: Hubs & Authorities

Proposed by Kleinberg. Query-dependent.


**Authorities:** Pages with good content (linked to by many hubs).

**Hubs:** Pages with good lists of links (link to many authorities).

Mutually reinforcing:

Good hubs point to good authorities.

Good authorities are pointed by good hubs.



# HITS Calculation

## Scenario

Graph: Hub H points to Auth A1 and Auth A2.

Goal: Find best Hubs/Auths.

Start: All scores = 1.

Iter 1: Update Authorities (Sum of Hubs)

$$\text{Auth(A1)} = \text{Hub(H)} = 1$$

$$\text{Auth(A2)} = \text{Hub(H)} = 1$$

Iter 1: Update Hubs (Sum of Auths)

$$\text{Hub(H)} = \text{Auth(A1)} + \text{Auth(A2)} = 1 + 1 = 2$$

Result: Hub H score grew because it points to valid Auths.

# The Power of Anchor Text

---



## Descriptive Power

Anchor text (the clickable text in a link) often provides a better description of the target page than the page itself.

- **Example:** Many people link to the IBM home page using the text "computer giant", even if the page itself only says "International Business Machines".
- **Google Bombs:** Historically used to manipulate rankings by mass-linking with specific phrases (e.g., "miserable failure").

# Quality Signals

## TrustRank

Start with a seed set of trusted pages (universities, news).

Propagate trust down the link graph.

Spam is usually far from trust.

## Freshness

How often to recrawl?

**News:** Minutes.

**Blogs:** Daily.

**Static:** Monthly.

Use adaptive crawl schedules based on change history.

# **Web Search – Part V:**

## **Spam Filter**

# The Spam Arms Race

If ranking = money, people will cheat.

## Content Spam

Keyword stuffing, hidden text, scraped content, auto-generated gibberish.

## Link Spam

Link farms, paid links, comment spam, buying expired domains.

## Cloaking

Showing search engines one version of a page and users another (e.g., porn/gambling).

# Generative vs. Discriminative

## Discriminative (Decision Boundary)

Logic: "Draw a line between Cats and Dogs."

Examples: KNN, SVM, Neural Networks.

Model:  $P(Y|X)$  directly.

## Generative (Probability Model)

Logic: "Learn what a Dog looks like. Learn what a Cat looks like. Compare new image to both."

Examples: Naive Bayes, LDA.

Model:  $P(X|Y)$  and  $P(Y)$ .

# Bayes' Theorem & Naive Assumption

$$P(\text{Class}|\text{Data}) = [P(\text{Data}|\text{Class}) * P(\text{Class})] / P(\text{Data})$$

## Key Terms

- Posterior:  $P(C|D)$  - What we want.
- Likelihood:  $P(D|C)$  - Prob. of data given class.
- Prior:  $P(C)$  - Baseline prob. of class.
- Evidence:  $P(D)$  - Constant (can ignore).

## The "Naive" Assumption

Features are independent given the class.

$$\left\{ \begin{array}{l} P(\text{Words} \mid \text{Spam}) \approx P(\text{Word1} \mid \text{Spam}) * \\ P(\text{Word2} \mid \text{Spam}) \dots \end{array} \right.$$

This turns a complex joint probability into simple multiplication.

# Spam Filtering: The Setup

---

## Training Data

We analyze existing emails to calculate probabilities.

Priors:

- $P(\text{Normal page}) = 0.625$
- $P(\text{Spam page}) = 0.375$

## Likelihoods $P(\text{Word}|\text{Class})$

Word	Normal	Spam
Moodle	0.38	0.12
Grade	0.30	0.25
Test	0.23	0.00
Money	0.07	0.63

# Classifying: "Money Grade"

Score for Normal (N)

$$\begin{aligned} P(N) * P(\text{Money}|N) * P(\text{Grade}|N) \\ = 0.625 * 0.07 * 0.30 \\ = 0.013 \end{aligned}$$

Score for Spam (S)

$$\begin{aligned} P(S) * P(\text{Money}|S) * P(\text{Grade}|S) \\ = 0.375 * 0.63 * 0.25 \\ = 0.059 \end{aligned}$$

Result: SPAM ( $0.059 > 0.013$ )

# Practical Issues & Solutions

## 1. Underflow

Multiplying many small probabilities results in computer zero.

Solution: Log Probabilities

$$\log(a \cdot b) = \log(a) + \log(b)$$

We sum logs instead of multiplying.

## 2. Zero Frequency

If a word (e.g., "Test") never appears in Spam, probability becomes 0, wiping out the whole score.

Solution: Smoothing (Laplace)

Add 1 to all counts so no probability is ever exactly zero.

# Summary

We have traversed the stack from the physical Infrastructure (IPs/DNS), through the Crawling layer (BFS/Politeness), modeled the Web Graph (SCC/Bow-Tie), applied Link Analysis (PageRank), and finally implemented Machine Learning (Naive Bayes) to classify content.